The Euclidean algorithm is not optimal among uniform processes

Jason Davies*

October 23, 2026

Abstract

Reusing a stored Euclidean prefix, we derive $\gcd(a, b)$ from a, b with $O(\log a / \log\log a)$ applications of the remainder function alone, uniformly for $a > b \geq 1$. These derivations yield a single uniform process with the same bound, refuting Moschovakis's conjectured logarithmic lower bounds for uniform processes. The Main Conjecture for algorithms remains open. We also prove an $\Omega(\log\log a)$ lower bound for almost all coprime pairs and give a family of finite nondeterministic gcd programs with arbitrarily small logarithmic cost coefficients. The arithmetic and program bounds are formalized in Lean 4.

AI transparency statement. This paper was produced entirely by artificial intelligence, including the identification of the main result, its formalization in Lean, and the writing of the manuscript. The human author initiated the broader research and directed the AI systems, but did not write the text. Details are given at the end of the paper.

1 Introduction

Van den Dries and Moschovakis asked whether the Euclidean algorithm is optimal among its peers, the algorithms that compute the gcd from the remainder function [8, p. 390]. Moschovakis's Main Conjecture asserts that for every algorithm α from rem and the zero test eq_0 computing $\gcd(x, y)$ there is an $r > 0$ such that α makes more than $r \log x$ remainder calls on infinitely many pairs $x > y \geq 1$ [6, p. 2]. We refute its formulations for uniform processes, the book's more general model (Section 5); for algorithms it remains open. Throughout, "the book" means Moschovakis's monograph [6], and logarithms are to base 2.

Derivations. A *derivation* from a, b is a finite list that starts with a, b and in which each further entry is $\mathrm{rem}(x, y)$ for earlier entries x, y with $y > 0$. Each further entry is one *instruction*, that is, one remainder call. Entries already in the list are *available* and may be reused; no constants and no other operations are supplied. Let $G(a, b)$ be the least number of instructions in a derivation containing $\gcd(a, b)$, and let $D(a, b) = G(a, b)$ for coprime inputs. Thus $G(a, 0) = 0$ and $D(a, 1) = 0$. Arithmetic used in the proofs, such as subtraction, sorting or comparison, only selects instructions; it is never itself an instruction.

Theorem 1. *Uniformly for integers $a > b \geq 1$,*

$$G(a, b) = O\left(\frac{\log a}{\log\log a}\right).$$

In particular this bound holds for $D(a, b)$ when $\gcd(a, b) = 1$.

*jason@jasondavies.com

Idea of the proof. For sufficiently large n, let $a > b$ be coprime with $a \le N = 2^n$. Store a Euclidean prefix $r_0 = a, r_1 = b, \dots, r_h$, where $h = \lceil \sqrt{n} \rceil$. If the prefix reaches a positive remainder at most $N^{3/4}$, take one further remainder and recurse on the resulting smaller pair. Otherwise its consecutive pairs stay distinct modulo every $m > N^{1/3}$, so the prefix occupies at least $\sqrt{h}$ residue classes modulo m (Lemma 3). A close pair of residues lets three remainder calls produce a positive value below $m/(\sqrt{h} - 1)$ (Lemma 2). Reusing the same prefix, we repeatedly shrink the modulus from r_h to at most $N^{1/3}$ in $O(n/\log n)$ calls. Two recursive gcd derivations on inputs at most $N^{1/3}$ finish. If $T(n)$ is the worst derivation cost for inputs at most 2^n, the resulting recurrence (5) gives $T(n) = O(n/\log n)$. Tuning the prefix length sharpens the upper coefficient to $6 + o(1)$ (Corollary 4).

Consequences and further results. Uniform processes are the book's abstract model of computation from given primitives; every deterministic or nondeterministic recursive program induces one [6, §4A, Proposition 4A.1]. Section 5 converts the derivations into small certificates and uses the book's best-uniform-process theorem to obtain a single process with the same asymptotic bound on every input. This refutes Conjecture 3.4 of [5] and version (4F-5) of the Main Conjecture, and limits what the homomorphism method can prove.

We also prove that $D(a, b) \ge (1 - o(1)) \log\log N$ for all but $o(N^2)$ coprime pairs $2 \le b < a \le N$ (Proposition 9). Finally, fixing the prefix length gives finite nondeterministic programs with arbitrarily small logarithmic coefficients, uniformly over all inputs (Theorem 10). Their program size grows as the requested coefficient decreases.

Related work. The known intrinsic lower bound is $\Omega(\log\log a)$ on special families of coprime pairs [8, 9], [6, Theorem 6C.5]. Pratt's unpublished nondeterministic algorithm [6, Theorem 2E.2 and Problems x2E.17–x2E.19] matches this order on successive Fibonacci numbers [6, Corollary 6C.7]. Our almost-all bound gives the same order on a set of full density among coprime pairs.

In a different model, Bshouty obtained the optimal bound $\Theta(\log N/\log\log N)$ for gcd on inputs in $[0, N]$ using a random access machine with linear operations, shifts and indirect addressing. His construction reads precomputed blocks of the binary gcd algorithm [7], [3, §4.5.2] from a table [1, Theorems i and ii]. Here the saving comes from reusing a Euclidean prefix with remainder as the only instruction.

2 Normalization

Write $g = \gcd(a, b) > 0$. For positive y,

$$\operatorname{rem}(gx, gy) = g \operatorname{rem}(x, y).$$

Every generated value is divisible by g. Consequently scaling a derivation, or dividing all its values by g, preserves its length and the validity of every instruction. Hence

$$G(a, b) = D(a/g, b/g). \tag{1}$$

Define

$$T(n) = \max_{0 \le b \le a \le 2^n,\ a > 0} G(a, b).$$

Euclid's remainders form a derivation containing $\gcd(a, b)$, so G and T are finite. By (1) it suffices to bound D; the cases $a = b$, $b = 0$ and $b = 1$ need no instructions.

3 Three-call contractions

Lemma 2. *If $m > 0$ is available and a finite collection S of available values occupies $t \geq 2$ residue classes modulo m, at most three instructions generate w with*

$$0 < w < \frac{m}{t-1}.$$

Proof. Sort the distinct residues. Some adjacent pair $u < v$ has gap $\delta = v - u < m/(t-1)$, because the sum of the $t-1$ adjacent gaps is less than m. Each selected residue is $\mathrm{rem}(x, m)$ for some $x \in S$, so it costs one instruction. If $u = 0$, use v. If $0 < u \leq \delta$, use u. Otherwise $0 < \delta < u$ and $v = u + \delta$, so $\mathrm{rem}(v, u) = \delta$. This uses at most three instructions. □

Lemma 3 (Euclidean-prefix lemma)**.** *Let $r_0 = a > r_1 = b > \cdots > r_h > 0$ be a Euclidean prefix of a coprime pair, with $r_{i-1} = q_i r_i + r_{i+1}$ and $q_i \geq 1$ for $1 \leq i < h$. For every integer $m > 2a/r_h$, the h pairs*

$$(r_i \bmod m, r_{i+1} \bmod m), \qquad 0 \leq i < h,$$

are distinct. The prefix therefore occupies t residue classes with $t^2 \geq h$.

Proof. Suppose the pairs at $i < j$ coincide. Set

$$M = \prod_{\ell=i+1}^{j} \begin{pmatrix} q_\ell & 1 \\ 1 & 0 \end{pmatrix} = \begin{pmatrix} A & B \\ C & E \end{pmatrix}, \qquad \begin{pmatrix} r_i \\ r_{i+1} \end{pmatrix} = M \begin{pmatrix} r_j \\ r_{j+1} \end{pmatrix}.$$

Right multiplication by $\left(\begin{smallmatrix} q & 1 \\ 1 & 0 \end{smallmatrix}\right)$ sends (A, B, C, E) to $(qA+B,\ A,\ qC+E,\ C)$, starting from $(q, 1, 1, 0)$ for a single factor. By induction, all entries are nonnegative, $A \geq B \geq E$, $A \geq C \geq 1$, $B \geq 1$, and $\det M = (-1)^{j-i}$. Also $Ar_j \leq r_i \leq a$ and $r_j > r_h$ (as $j < h$), whence $A < a/r_h$. Write

$$\Delta = \det(M - I) = 1 + \det M - A - E.$$

For an odd number of factors, $\Delta = -A - E < 0$. By the same recursion, a product of at least two factors has $A \geq 2$ and $E \geq 1$; thus for an even positive number of factors, $\Delta = 2 - A - E < 0$. In either case,

$$0 < |\Delta| \leq A + E \leq 2A \leq 2a/r_h < m.$$

The presumed congruence gives $(M - I)(r_j, r_{j+1})^{\mathsf{T}} \equiv 0 \pmod{m}$. Multiplication by the adjugate shows that m divides both Δr_j and Δr_{j+1}. Consecutive remainders are coprime, so Bézout's identity implies $m \mid \Delta$, a contradiction. Finally, t individual residue classes give at most t^2 ordered pairs. □

Since Lemma 3 holds for every $m > 2a/r_h$, one stored prefix serves every later modulus, although these moduli are not known in advance. When $\sqrt{h} > 1$, Lemma 2 applied to the prefix gives three instructions that generate

$$0 < w < \frac{m}{\sqrt{h} - 1}. \tag{2}$$

We call these three instructions a *contraction* of m, and w the new modulus.

4 Construction and recurrence

Fix an integer $n > 16$ and put $N = 2^n$ and $h = \lceil \sqrt{n} \rceil > 4$. After normalization, assume $a > b > 1$ are coprime and $a \leq N$. We use $N^{3/4}$ as the *early threshold* and $N^{1/3}$ as the *target modulus*. Any thresholds N^α, N^β with $\alpha < 1$, $\alpha + \beta > 1$ and $2\beta < 1$ would do: these make the early recursive problem smaller, let the prefix lemma apply above the target, and make the two final problems together smaller. We take $\alpha = 3/4$ and $\beta = 1/3$.

Phase (i): the prefix. Starting from $r_1 = b$, generate the Euclidean remainders $r_2, r_3, \ldots$, one instruction each, until the most recent value is at most $N^{3/4}$ or r_h is available. This costs at most h instructions. Since the inputs are coprime, Euclid's algorithm reaches 1 before 0, so the stopping value is positive.

If the stopping value c is at most $N^{3/4}$, generate the next remainder $s < c$ with one more instruction. The pair (c, s) is coprime, with both entries at most $2^{\lceil 3n/4 \rceil}$. This gives the bound

$$h + 1 + T(\lceil 3n/4 \rceil). \tag{3}$$

Phase (ii): contractions. Otherwise $r_h > N^{3/4}$, so $2a/r_h < 2N^{1/4} < N^{1/3}$ (as $n > 12$), and Lemma 3 applies to every available $m > N^{1/3}$. Starting at $m = r_h$, repeatedly replace m by the value w of (2) until $0 < m \le N^{1/3}$. Set

$$\kappa = \sqrt{h} - 1 > 1, \qquad J = \left\lceil \frac{2n/3}{\log \kappa} \right\rceil.$$

There are at most J contractions: if J have occurred then $m < N/\kappa^J \le N^{1/3}$. Since $\log \kappa = \frac{1}{4} \log n + O(1)$, their cost is $3J = O(n/\log n)$.

Phase (iii): two recursive derivations. The small m need not be coprime to either input. Generate $u = \mathrm{rem}(a, m)$ and $v = \mathrm{rem}(b, m)$, then append a derivation of $d = \gcd(m, u)$ from m, u and one of $e = \gcd(d, v)$ from d, v. Since $m \le 2^{\lceil n/3 \rceil}$ and G is symmetric, each costs at most $T(\lceil n/3 \rceil)$. Finally

$$e = \gcd(\gcd(m, \mathrm{rem}(a, m)), \mathrm{rem}(b, m)) = \gcd(m, \gcd(a, b)) = 1.$$

The cost in this case is at most

$$h + 3J + 2 + 2T(\lceil n/3 \rceil). \tag{4}$$

The recurrence. The two case estimates (3) and (4) and normalization imply

$$T(n) \le \max\{T(\lceil 3n/4 \rceil) + O(\sqrt{n}),\ 2T(\lceil n/3 \rceil) + O(n/\log n)\}. \tag{5}$$

All implied constants are independent of the inputs.

Proof of Theorem 1. For $n \ge 2$ let $f(n) = n/\log n$. The ratios

$$\frac{f(\lceil 3n/4 \rceil)}{f(n)} \longrightarrow \frac{3}{4}, \qquad \frac{2f(\lceil n/3 \rceil)}{f(n)} \longrightarrow \frac{2}{3}$$

show that, beyond a fixed threshold, $f(\lceil 3n/4 \rceil) \le \frac{7}{8} f(n)$ and $2f(\lceil n/3 \rceil) \le \frac{5}{6} f(n)$. Both overheads in (5) are at most $c_0 f(n)$ for a fixed c_0. Choose the induction threshold above 16, where the recurrence holds, and large enough that both recursive indices are at least 2 and strictly less than n. Choose $C \ge 8c_0$, increased to cover the finitely many values below the threshold. Strong induction then gives $T(n) \le Cf(n)$: the two bounds are $(7C/8 + c_0) f(n)$ and $(5C/6 + c_0) f(n)$, both at most $Cf(n)$. Finally take $n = \lceil \log a \rceil$. □

4.1 The leading coefficient

The same construction allows a more precise bound without changing either arithmetic lemma.

Corollary 4. *For every $\varepsilon > 0$, there is an a_0 such that for every $a \geq a_0$ and $0 \leq b \leq a$,*

$$G(a,b) \leq (6+\varepsilon)\frac{\log a}{\log\log a}.$$

Equivalently, the uniform upper bound has coefficient $6 + o(1)$.

Proof. Fix $0 < c < 1/2$ and $0 < \beta < 1/2$, and put $\alpha = 1 - \beta/2$. For large n put

$$p = \lfloor c \log n \rfloor, \qquad k = 2^p, \qquad h = k^2 + 1, \qquad J = \lceil n/p \rceil.$$

Use N^α as the early threshold and N^β as the target modulus, where $N = 2^n$. In the late case $r_h > N^\alpha$, so $2a/r_h < 2N^{1-\alpha} = 2N^{\beta/2} < N^\beta$ for sufficiently large n, and the prefix lemma applies to every modulus above N^β. Since $k^2 < h$, the prefix occupies at least $k+1$ residue classes, so each contraction reduces the modulus by a factor greater than k. Moreover $k^J \geq N$, so J contractions suffice. Writing $f(n) = n/\log n$, we have

$$h = O(n^{2c}) = o(f(n)), \qquad 3J = \left(\frac{3}{c} + o(1)\right) f(n).$$

The two branches therefore give

$$T(n) \leq \max\left\{ T(\lceil \alpha n \rceil) + o(f(n)), \quad 2T(\lceil \beta n \rceil) + \left(\frac{3}{c} + o(1)\right) f(n) \right\}.$$

Theorem 1 implies that $\Lambda = \limsup_{n\to\infty} T(n)/f(n)$ is finite and nonnegative. For each fixed $q > 0$, the ratio $f(\lceil qn \rceil)/f(n)$ tends to q. Taking upper limits in the recurrence yields

$$\Lambda \leq \max\{\alpha\Lambda,\ 2\beta\Lambda + 3/c\}.$$

Since $\alpha < 1$ and $\Lambda \geq 0$, this implies $(1 - 2\beta)\Lambda \leq 3/c$. First let $\beta \downarrow 0$ and then $c \uparrow 1/2$ to obtain $\Lambda \leq 6$. Finally, $n = \lceil \log a \rceil$ satisfies $f(n) \sim \log a/\log\log a$, proving the assertion. All upper limits concern the worst-case function T, so the resulting threshold is independent of b. □

We do not claim that 6 is optimal.

5 Certificates and uniform processes

Structures and certificates. Write tt, ff for the truth values and eq_c for the test $x = c$. We use the book's partial remainder operation, defined only when its second argument is nonzero, in the structures

$$\mathcal{N}_0 = (\mathbb{N}, \mathrm{rem}, \mathrm{eq}_0), \qquad \mathcal{N}_\varepsilon = (\mathbb{N}, \mathrm{rem}, \mathrm{eq}_0, \mathrm{eq}_1).$$

The latter is the book's Euclidean structure $\mathbf{N}_\varepsilon$ [6, (1C-10), p. 24]; our $\mathcal{N}_0$ should not be confused with the book's $\mathbf{N}_0$, which has no remainder operation [6, p. 171]. Coprimeness is the relation $x \perp\!\!\!\perp y \iff x, y \geq 1 \wedge \gcd(x,y) = 1$ [6, §1A].

A derivation records finitely many remainder equations. These form a partial structure generated by the inputs. A *certificate* for f at $\vec{x}$ in a structure $\mathcal{A}$ is such a finite generated partial substructure, possibly with test facts, on which every homomorphism into $\mathcal{A}$ respects f at $\vec{x}$. In other words, any assignment of values preserving the recorded operations and tests must preserve the answer. The *intrinsic remainder-call complexity* $\mathrm{calls}_{\mathrm{rem}}(\mathcal{A}, f, \vec{x})$ is the least number of remainder facts in a certificate [6, §§1D, 4C–4E].

Proposition 5. *For coprime $a, b \geq 1$, $\mathrm{calls}_{\mathrm{rem}}(\mathcal{N}_\varepsilon, \perp\!\!\!\perp, a, b) = D(a, b)$. For $(a, b) \neq (0, 0)$,*

$$G(a, b) \leq \mathrm{calls}_{\mathrm{rem}}(\mathcal{N}_0, \gcd, a, b) \leq G(a, b) + 2,$$

and the same bounds hold for $\mathrm{calls}_{\mathrm{rem}}(\mathcal{N}_\varepsilon, \gcd, a, b)$.

Proof. A derivation of 1, together with $\mathrm{eq}_1(1) = \mathrm{tt}$, $\mathrm{eq}_0(a) = \mathrm{ff}$ and $\mathrm{eq}_0(b) = \mathrm{ff}$, certifies $\perp\!\!\!\perp$: under any homomorphism the input images are positive, and each of their common divisors divides every transported generated value, hence divides the image 1. Conversely, if a certificate has no positive eq_1 fact, then $x \mapsto 2x$ preserves its remainder facts, zero tests and negative one-tests, but not coprimeness. So every certificate contains 1, and listing its values in generation order gives a derivation of 1 with at most as many instructions as the certificate has remainder facts. This proves the first assertion.

For gcd as a function, the certificate must reach its output g, giving a lower bound of $G(a, b)$. For an upper bound, append

$$\mathrm{rem}(a, g) = z, \qquad \mathrm{rem}(b, g) = z, \qquad \mathrm{eq}_0(z) = \mathrm{tt},$$

where the generated z is zero. Every homomorphism π has $\pi(z) = 0$ and $\pi(g) > 0$ (by definedness of either appended remainder). Thus $\pi(g)$ divides both input images. Conversely every common divisor of those images divides $\pi(g)$ by induction over the derivation. This proves the bounds for $\mathcal{N}_0$. The lower-bound argument applies to every gcd certificate, eq_1 facts included, and allowing eq_1 facts can only lower the count, so the upper bound transfers to $\mathrm{calls}_{\mathrm{rem}}(\mathcal{N}_\varepsilon, \gcd, a, b)$. □

Remark 6. At $(0, 0)$ the output is an input and the empty certificate suffices. For gcd restricted to positive inputs, appending $\mathrm{eq}_0(a) = \mathrm{ff}$ and $\mathrm{eq}_0(b) = \mathrm{ff}$ keeps input images in the domain, at no remainder cost. For a noncoprime input, adding $\mathrm{eq}_1(g) = \mathrm{ff}$ to the gcd certificate certifies noncoprimeness, so the sublogarithmic remainder-call bound holds for both truth values of coprimeness as well as for the gcd. The overhead is a fixed additive constant and does not change the leading coefficient, so Corollary 4 holds for all three complexities of Corollary 8 below on positive inputs.

A single uniform process. Proposition 5, Remark 6, and Theorem 1 bound each intrinsic complexity by $O(\log a / \log\log a)$, uniformly for $a \geq b \geq 1$. The book's Theorem 4F.1 states that, whenever every input has a certificate, one uniform process attains the intrinsic complexity at every input [6, pp. 149–150]. Hence:

Corollary 7. *For each of gcd and coprimeness, there is a single uniform process of $\mathcal{N}_\varepsilon$ computing it on $\mathbb{N}^2$ with $O(\log a / \log\log a)$ remainder calls, uniformly for $a \geq b \geq 1$. The gcd process can also be taken over $\mathcal{N}_0$.*

Proof. The certificate bounds give the cost assertion by Theorem 4F.1. Every input has a certificate: symmetry covers reversed inputs, and a zero-test fact suffices when either input is zero. Thus the same process is defined on all of $\mathbb{N}^2$. □

The corollary rules out the possibility raised after Corollary 6C.7 of the book that the Main Conjecture holds for all uniform processes of $\mathcal{N}_\varepsilon$ [6, p. 188].

The conjectures refuted. Moschovakis's 2012 notes state:

> **Conjecture 3.4** [5, p. 19]. There is a rational number $r > 0$ such that for infinitely many $a \geq b \geq 1$, $\mathrm{calls}_{\mathrm{rem}}(\mathcal{N}_\varepsilon, \perp\!\!\!\perp, a, b) \geq r \log a$.

The book formulates one precise version of the Main Conjecture [6, p. 151]: there are a set $\mathcal{D} \subseteq \mathbb{N}^2$ and a $K > 0$ such that

$$(x, y) \in \mathcal{D} \implies \mathrm{calls}_{\mathrm{rem}}(\mathrm{Eucl}, x, y) \leq L(x, y) \leq K \,\mathrm{calls}_{\mathrm{rem}}(\mathcal{A}, \gcd, x, y), \tag{4F-5}$$

where $L(x, y) = 2 \log \max(x, y)$ and $\mathrm{calls}_{\mathrm{rem}}(\mathrm{Eucl}, x, y)$ counts the remainder calls of the Euclidean algorithm Eucl (the book writes ε) on (x, y). The book states (4F-5) for the Euclidean algorithm in $\mathbf{N}_\varepsilon$, so its $\mathcal{A}$ is $\mathcal{N}_\varepsilon$; we treat both $\mathcal{A} = \mathcal{N}_\varepsilon$ and $\mathcal{A} = \mathcal{N}_0$. The book adds that $\mathcal{D}$ then satisfies its conditions (4F-3) and (4F-4); the latter makes the intrinsic complexity unbounded on $\mathcal{D}$, so $\mathcal{D}$ is infinite (for finite $\mathcal{D}$ the statement is trivial). We refute the weaker statement that drops the first inequality and asks only that $\mathcal{D}$ be infinite.

Corollary 8. *The intrinsic complexities*

$$\mathrm{calls}_{\mathrm{rem}}(\mathcal{N}_\varepsilon, \perp\!\!\!\perp, a, b), \qquad \mathrm{calls}_{\mathrm{rem}}(\mathcal{N}_0, \gcd, a, b), \qquad \mathrm{calls}_{\mathrm{rem}}(\mathcal{N}_\varepsilon, \gcd, a, b)$$

are $O(\log a / \log \log a)$, uniformly for $a \geq b \geq 1$. Consequently Conjecture 3.4 of [5] is false, and so is the version (4F-5) of the Main Conjecture with $\mathcal{A} = \mathcal{N}_0$ or $\mathcal{A} = \mathcal{N}_\varepsilon$: for every $r > 0$ only finitely many pairs $a \geq b \geq 1$ satisfy $\mathrm{calls}_{\mathrm{rem}}(\mathcal{N}_\varepsilon, \perp\!\!\!\perp, a, b) \geq r \log a$, and for every $K > 0$ only finitely many pairs $(x, y) \in \mathbb{N}^2$ satisfy $L(x, y) \leq K \,\mathrm{calls}_{\mathrm{rem}}(\mathcal{A}, \gcd, x, y)$.

The corollary follows from Theorem 1 and Proposition 5; for noncoprime pairs, the bound on the coprimeness complexity comes from Remark 6. For (4F-5), pairs with $x < y$ or a zero entry are handled as in the proof of Corollary 7; the latter need no remainder facts.

The limit of the homomorphism method. The Homomorphism Test proves lower bounds on intrinsic complexity by showing that every generated diagram below the proposed cost admits a homomorphism that fails to preserve the answer [6, Lemma 4E.3]. Such bounds apply to every uniform process computing the function, hence to every recursive program [6, Proposition 4E.1, Corollary 4E.2]. By Corollary 8, no lower bound on the intrinsic complexities listed there can grow faster than $\log a / \log \log a$, so a logarithmic lower bound for algorithms must use restrictions beyond uniformity, such as finite program syntax or determinism.

6 An almost-all lower bound

The difficult pairs of [6, Theorem 6C.5] give $\Omega(\log \log a)$ on a sparse set. An elementary Diophantine argument gives the same order for almost all coprime pairs.

Proposition 9. *For integers $d \geq 1$ and $N \geq 3$, at most $40(2 + \ln N) N^{2 - 2^{-d}}$ coprime pairs $2 \leq b < a \leq N$ satisfy $D(a, b) \leq d$. Hence $D(a, b) \geq (1 - o(1)) \log \log N$ for all but $o(N^2)$ of these pairs.*

Proof. For $B \geq 1$ call (a, b) *B-regular* if $|qb - pa| \geq a/(Bq)$ whenever $q \in \mathbb{N}$, $q \geq 1$, $p \in \mathbb{Z}$ and $qb \neq pa$. We first bound the cost on regular pairs, then count the irregular ones.

Every generated value has an integer representation $v = ua + wb$. On a regular pair, $v > 0$ implies $v \geq a/(B(|u| + |w|))$: for $w \neq 0$ use regularity with $q = |w|$, and for $w = 0$ note $v \geq a$. Let H_t bound the coefficient norm $|u| + |w|$ in one such representation of each value available after

t instructions, starting with $H_0 = 1$. In an instruction $r = \mathrm{rem}(x, y)$, write $x = sy + r$. Since generated values are at most a, the quotient satisfies $s \le a/y \le BH_t$. Subtracting s times the representation of y from that of x gives

$$H_{t+1} \le H_t + BH_t^2 \le 2BH_t^2, \qquad H_t \le (2B)^{2^t-1}.$$

Producing 1 within d instructions therefore forces

$$1 \ge \frac{a}{BH_d}, \qquad a \le \frac{(2B)^{2^d}}{2}.$$

To count irregular pairs below N, choose a witness $0 < |qb - pa| < a/(Bq)$. Dividing p, q by their gcd preserves this inequality, so assume $\gcd(p, q) = 1$. The nonzero integer $r = qb - pa$ then gives $q < a/B \le N/B$ and $0 \le p \le q$. For each q, there are at most $2q$ choices of p and $2N/(Bq)$ choices of the signed integer r. Given p, q, r, the equation $pa \equiv -r \pmod q$ determines one residue class for a, leaving at most $N/q + 1$ choices; b is then determined. With $Q = \lfloor N/B \rfloor$, the number of irregular pairs is at most

$$\begin{aligned}\sum_{q=1}^{Q} (2q)\frac{2N}{Bq}\left(\frac{N}{q} + 1\right) &= \frac{4N^2}{B}\sum_{q=1}^{Q}\frac{1}{q} + \frac{4NQ}{B} \\ &\le \frac{4N^2}{B}(2 + \ln N),\end{aligned}$$

using $Q \le N$, $B \ge 1$, and the harmonic-sum bound; an empty sum is zero.

Now split into dyadic ranges $X/2 < a \le X$, and take $B = X^{2^{-d}}/2$. If $B \ge 1$, a regular pair with $D \le d$ would have $a \le X/2$, so all such pairs in the range are irregular. Their number is at most $8(2 + \ln X)X^{2-2^{-d}}$. When $B < 1$, the same bound follows from the trivial bound X^2, since $X^{2^{-d}} < 2$. Put $s = 2 - 2^{-d} \ge 3/2$. Summing over the nonempty ranges with $X = N/2^j \ge 3$, and bounding their logarithmic factors by $2 + \ln N$, gives

$$8(2 + \ln N)N^s \sum_{j \ge 0} 2^{-js} = \frac{8}{1 - 2^{-s}}(2 + \ln N)N^s \le 40(2 + \ln N)N^s.$$

For the second assertion, take $d = \lfloor \log\ln N - 2\log\ln\ln N \rfloor$ for sufficiently large N. Then $N^{-2^{-d}} \le \exp(-(\ln\ln N)^2)$, so the exceptional set has size $o(N^2)$, while $d = (1 - o(1))\log\log N$. □

The saving 2^{-d} in the exponent of $N^{2-2^{-d}}$ shrinks as d grows, so this argument does not reach the upper-bound scale $\log N/\log\log N$.

7 Fixed programs with arbitrarily small coefficients

For a nondeterministic program P, let $c_P(a, b)$ be the minimum number of remainder calls among its convergent computations. We require at least one such computation on each input and require every convergent computation to return the gcd. These are the conventions of [6, (2E-2), (2E-5)].

Theorem 10. *For each integer $p \ge 1$, there is a finite nondeterministic recursive program P_p over $\mathcal{N}_0$ computing the gcd such that*

$$\limsup_{a\to\infty} \max_{1 \le b \le a} \frac{c_{P_p}(a, b)}{\log a} \le \frac{3}{p}.$$

Its generator, defined below, has $W = 4^p + 5$ individual parameters, and no internal function has larger arity. Consequently, for every $\varepsilon > 0$ there are a fixed program P and a constant $B \geq 0$ such that

$$\forall a \geq b \geq 1, \qquad c_P(a, b) \leq \varepsilon \log a + B.$$

The program depends on ε but not on the inputs. Since the Main Conjecture lets its coefficient r depend on the program, the family is compatible with it: no single sublogarithmic program is obtained.

The program. Fix $W \geq 3$ and write $\vec{x} = (x_0, \ldots, x_{W-1})$ and $\vec{x}[i \leftarrow E]$ for the tuple with its ith expression replaced by E. The generator can return a parameter, replace one by a remainder, run a recursive subcomputation while retaining the caller's other values, or copy parameters into new positions. These four actions are expressed by the finite alternatives

$$\begin{aligned}
\mathrm{choose}(\vec{x}) &= x_i && (0 \leq i < W), && \text{(C)}\\
\mathrm{gen}(\vec{x}) &= \mathrm{choose}(\vec{x}), && && \text{(G1)}\\
\mathrm{gen}(\vec{x}) &= \mathrm{gen}(\vec{x}[i \leftarrow \mathrm{rem}(\mathrm{choose}(\vec{x}), \mathrm{choose}(\vec{x}))]) && (0 \leq i < W), && \text{(G2)}\\
\mathrm{gen}(\vec{x}) &= \mathrm{gen}(\vec{x}[i \leftarrow \mathrm{gen}(\vec{x})]) && (0 \leq i < W), && \text{(G3)}\\
\mathrm{gen}(\vec{x}) &= \mathrm{gen}(\mathrm{choose}(\vec{x}), \ldots, \mathrm{choose}(\vec{x})). && && \text{(G4)}
\end{aligned}$$

There are W independent selector calls in (G4). Tuple notation and indices are syntax abbreviations: after fixing W, each line expands into ordinary first-order terms with fixed arities. No tuple, index, numeric constant, or update operation is supplied as runtime data or a primitive. The nested generator in (G3) is evaluated once; its caller's other parameters remain available when it returns. The recursion stack can grow, so W is an arity bound, not a total storage bound.

Add $\mathrm{loop}(z) = \mathrm{loop}(z)$ and define $\mathrm{accept}(a, b, g)$ by nested zero-test conditionals: if $g = 0$, loop; otherwise test $\mathrm{rem}(a, g) = 0$ and then $\mathrm{rem}(b, g) = 0$, returning g if both tests succeed and looping otherwise. The head is

$$P_W(a, b) = \text{if } \mathrm{eq}_0(b) \text{ then } a \text{ else } \mathrm{accept}(a, b, \mathrm{gen}(a, b, a, \ldots, a)).$$

The four internal symbols have arities $W, W, 3, 1$ (for choose, gen, accept, loop). Every common divisor of the generator's parameters divides its output, by induction on a finite evaluation. Thus a candidate is divisible by $\gcd(a, b)$; successful verification proves the reverse divisibility. Every convergent computation therefore returns the gcd. There is always a convergent computation following Euclid, using parameter selection and copying to retain the current pair. For $b > 0$, final verification costs exactly two remainder calls.

This syntax is a finite recursive program in the sense of [6, §§2A, 2E]; Appendix A records the Lean check that its evaluation trees correspond, with equal remainder cost, to convergent computations of a specialization of the book's recursive machine.

Proof of Theorem 10. We follow the construction of Section 4: store the prefix, contract, then recurse twice through nested generator calls. Take $P_p = P_W$ with $W = 4^p + 5$; that is, set $k = 2^p$, $h = k^2 + 1$ (the prefix length, as in Corollary 4 but now with p fixed), and $W = h + 4$. Store $r_0, \ldots, r_h$ in $h + 1$ parameters, reserving three parameters for the modulus and two scratch residues. Lemmas 2 and 3 give a contraction by a factor greater than k in at most three remainder calls, preserving the prefix. The program can choose precisely the operands of each contraction.

The *generator cost* of (a, b) is the minimum number of remainder calls in an evaluation of gen that returns $\gcd(a, b)$ from $(a, b, a, \ldots, a)$. Let $T_W(n)$ be its maximum over $a, b \leq 2^n$. Euclid gives

$T_W(n) = O(n)$. To *route* to x, y means to apply (G4) to pass to the tuple $(x, y, x, \ldots, x)$, at no remainder cost. Copying parameters gives symmetry. Multiplying every value by $g = \gcd(a, b)$ turns a run on $(a/g, b/g)$ that returns 1 into a run on (a, b) that returns g, with the same remainder count. It therefore suffices to analyze coprime inputs.

Fix $0 < \beta < 1/2$ and put $\alpha = 1 - \beta/2$. Use the early threshold N^α and the target modulus N^β, where $N = 2^n$. If no prefix remainder is at most N^α, then $r_h > N^\alpha$ and $2a/r_h < 2N^{1-\alpha} < N^\beta$ for all sufficiently large n. There are at least $k+1$ distinct prefix residues modulo every modulus used before stopping, since $h > k^2$. At most $\lceil n/p \rceil$ contractions produce a positive $m \le N^\beta$.

In the early case, route directly to the smaller consecutive Euclidean pair. Otherwise form $u = a \bmod m$ and $v = b \bmod m$ in the scratch positions. A nested generator call, with its parameters routed to copies of m, u, returns $d = \gcd(m, u)$ while its caller retains v. Replace the modulus parameter by d and route to copies of d, v for the second gcd computation. The result is $\gcd(m, a, b) = 1$. All these routes and nested calls are available in the same fixed program. Thus, eventually in n,

$$T_W(n) \le \max\left\{T_W(\lceil \alpha n \rceil) + h + 1, \quad 2T_W(\lceil \beta n \rceil) + h + 3\lceil n/p \rceil + 2\right\}.$$

Set $\Lambda_W = \limsup_{n\to\infty} T_W(n)/n$, finite and nonnegative by Euclid. Since h is fixed, division by n and passage to upper limits gives

$$\Lambda_W \le \max\{\alpha\Lambda_W,\ 2\beta\Lambda_W + 3/p\}.$$

As $\alpha < 1$ and $\Lambda_W \ge 0$, this implies $(1 - 2\beta)\Lambda_W \le 3/p$. The program is independent of β, so letting $\beta \downarrow 0$ gives $\Lambda_W \le 3/p$. Take $n = \lceil \log a \rceil$ and add the two verification calls. Finally, choose $p > 3/\varepsilon$ and absorb the finitely many smaller inputs in B. □

8 Discussion and open problems

The gap. The intrinsic complexity of coprimeness in $\mathcal{N}_\varepsilon$ is at least $\Omega(\log\log a)$ on almost all coprime pairs and at most $O(\log a/\log\log a)$ on all positive pairs. Closing this gap is open; unlike Bshouty's bound in his model, our upper bound is not known to be optimal.

From a growing prefix to one program. The remaining obstacle in turning the construction into a sublogarithmic program is reusable access to a prefix whose length grows with the input. Selecting a close pair of residues costs nothing in a derivation, whereas computing all prefix residues would take about h calls per contraction. The programs of Section 7 select operands nondeterministically from a fixed number of parameters. Can one fixed first-order recursive program retain and repeatedly access a growing prefix at the same remainder cost? Proposition 2E.1 of the book [6, §2E] converts each certificate into a computation of some program, not of one program for all certificates.

Stronger lower-bound methods. The limit established in Section 5 leaves room for arguments combining homomorphisms with additional restrictions on computations. In particular, the conjecture of a logarithmic lower bound for deterministic uniform processes remains open [5, Conjecture 3.9, p. 23], [6, §4H and p. 188].

9 Computational checks

The offline Python planner `remainder_only_certificate.py` in the accompanying repository https://github.com/jasondavies/euclid-uniform-processes uses the parameters of Section 4 and unrestricted arithmetic to choose instructions. An independent replay then checks operand availability, nonzero divisors, every remainder equation and the final gcd. Below a base size (256 bits by default) it uses Euclid's algorithm. On one pseudorandom coprime pair of each size, generated from its default seed, it gives these instruction counts:

Input bits	This construction	Euclid (to reach 1)
512	196	326
2048	669	1204
8192	2059	4754

The asymptotic claim rests on the proof, not on these examples. The accompanying repository also checks scaled and noncoprime inputs, rejects invalid traces, and exhaustively tests the prefix lemma over a finite range. For Section 7, a separate compiler and interpreter replay successful choices in finite first-order programs, evaluating arguments right to left as the book's machine does. The `README.md` records commands, test ranges, seeds, and results. The Python interpreter is not formally verified; unrestricted arithmetic used to plan choices is outside the programs.

A Formal verification

The development uses Lean 4.35.0 [2] with mathlib [4] at its v4.35.0 tag; the pinned mathlib commit is `56a005855ea17e0b0958f73ec801ae140b0a1679`. The proved library contains no proof placeholders or added axioms; the main-result audit reports only `propext`, `Classical.choice`, and `Quot.sound`. Table 1 gives the main declarations; the repository README gives the complete map and reproduction commands. The general theory of uniform processes, including Theorem 4F.1 used in Corollary 7, and the historical lower bounds are cited, not formalized. The certificate bounds and conjecture refutations are formalized independently of that theory.

The development is also packaged for the Palomar registry: `Challenge.lean` imports only mathlib and states the main theorems with every definition visible and deliberate `sorry` placeholders. `Solution.lean` imports the proved library, independently of these placeholders. Palomar's comparator checks that the proved statements are exactly those of `Challenge.lean` and replays the proofs through Lean's kernel and independent kernels. The eleven compared theorems are marked † in Table 1.

Arithmetic bounds. Costs are lengths of explicit remainder derivations. The finite theorem `sizeCost_step` proves the recurrence from natural-number size parameters; it does not assume a recurrence for the cost function. For the initial Big-O proof, Lean uses $j = \lfloor \log n \rfloor$, $p = \lfloor j/4 \rfloor$ and $h = 2^{2p} + 1$ in place of $h = \lceil \sqrt{n} \rceil$, and checks the integer rounding estimates before passing to real logarithms. The refinement uses the tunable parameters of Corollary 4 and checks the rounding limits preserving coefficient 6. For Proposition 9, Lean takes $B = a^{2^{-d}}/2$ separately for each pair instead of splitting into dyadic ranges. It proves the stated constant 40 and the exceptional-set limit for every fixed $\varepsilon > 0$, with threshold $(1 - \varepsilon) \log \log N$.

Table 1: Main results and Lean declarations, all in the namespace `ShortGCD`. Names marked † are compared against `Challenge.lean`.

Paper	Lean declaration(s)
Theorem 1	`uniform_real_bound`, `sizeCost_isBigO`
Normalization (1)	`G_normalize`
Lemma 2	`small_remainder`
Lemma 3	`prefix_pairs_distinct`, `residue_card_sq`
Corollary 4	`uniform_real_bound_six`†, `uniform_intrinsic_bound_six`†
Proposition 5	`callsRem_eq_D`, `G_eq_D`, `callsCoprimeness_eq_callsRem`, `gcd_certificate_lower`, `callsGcd_le`, `callsRemGcd_le_callsGcd`, `G_le_callsRemGcd`†, `G_le_callsGcd`†, `callsCoprimeness_eq_G`†
Remark 6	`callsCoprimeness_le`
Corollary 7	Cited Theorem 4F.1; not formalized here
Corollary 8	`not_conjecture_3_4`†, `not_main_conjecture_4F5`†, `not_main_conjecture_4F5_Neps`†; literal versions `not_conjecture_3_4_literal`†, `not_main_conjecture_4F5_literal`†
Proposition 9	`AlmostAll.card_D_le`, `AlmostAll.tendsto_card_within_loglog`, `tendsto_card_callsCoprimeness_loglog`†
Theorem 10: cost	`RegisterProgram.maxCost_step`, `RegisterProgram.Fixed.limsup_profile_le`
Theorem 10: syntax	`RegisterProgram.Syntax.eval_head_correct`, `RegisterProgram.Syntax.eval_head_exists`, `RegisterProgram.Syntax.arbitrarily_small_program`
Operational correspondence	`RegisterProgram.Syntax.Machine.eval_iff_steps`, `RegisterProgram.Syntax.Machine.arbitrarily_small_program`

Programs and machine computations. Generator runs translate into evaluations of finite first-order syntax with exactly the same remainder count, including final verification. Every terminating evaluation is correct, and a successful evaluation exists on every input. Lean also proves exact-cost equivalence with a specialization of the book's two-stack machine: arguments are evaluated right to left, zero tests and branches are separate transitions, and only remainder transitions are charged [6, Table 1, pp. 75–76]. The forward proof embeds each argument evaluation in its calling context; the reverse proof decomposes a convergent computation into its argument and body evaluations. Closures represent substitution in machine states, not extra program data. This checks the language and program family used here; the repository README describes the Python compiler and interpreter and their relation to the Lean syntax.

Declaration of AI-generated research and authorship

This manuscript is the product of an AI-driven research process. During a broader investigation initiated by the human author, an AI system identified the principal result. The author then instructed it to formalize the result in Lean and to prepare this manuscript. The AI systems developed the result and its presentation, including the mathematical exposition, the formalization, and the organization and writing of the text; the human author wrote no part of the manuscript. The research, the Lean development and the manuscript were produced with GPT-6-Astra (OpenAI), used through Codex. Claude Opus 5.5 (Anthropic), used through Claude Code, carried out review and editing passes and packaged the formalization for the Palomar registry, which included writing `Challenge.lean`, the lower-bound module and a proof change needed by the independent kernels.

The Lean development gives a machine-checkable version of the main claims, subject to its trusted foundations (Lean's kernel, mathlib, and the axioms listed in Appendix A) and to the correspondence between its definitions and the book's notions discussed there.

This declaration distinguishes the provenance of the work from conventional AI-assisted writing. The human author accepts responsibility for the submission and for the accuracy of this disclosure.

References

[1] Nader H. Bshouty. *Euclidean GCD algorithm is not optimal*. Research Report 1989-376-38, Department of Computer Science, University of Calgary, December 1989. doi:10.11575/PRISM/30475.

[2] Leonardo de Moura and Sebastian Ullrich. The Lean 4 theorem prover and programming language. In *Automated Deduction – CADE 28*, Lecture Notes in Computer Science, vol. 12699, pp. 625–635. Springer, 2021. doi:10.1007/978-3-030-79876-5_37.

[3] Donald E. Knuth. *The Art of Computer Programming, Vol. 2: Seminumerical Algorithms*, 3rd ed. Addison-Wesley, Reading, MA, 1997.

[4] The mathlib Community. The Lean mathematical library. In *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs (CPP 2020)*, pp. 367–381. ACM, 2020. doi:10.1145/3372885.3373824.

[5] Yiannis N. Moschovakis. *Absolute lower bounds for uniform processes*. Unpublished lecture notes (first draft), January 25, 2012. `https://www.math.ucla.edu/~ynm/nadalb.pdf`.

[6] Yiannis N. Moschovakis. *Abstract Recursion and Intrinsic Complexity*. Lecture Notes in Logic, vol. 48. Cambridge University Press and Association for Symbolic Logic, 2019. doi:10.1017/9781108234238.

[7] J. Stein. Computational problems associated with Racah algebra. *Journal of Computational Physics* 1(3):397–405, 1967. doi:10.1016/0021-9991(67)90047-2.

[8] Lou van den Dries and Yiannis N. Moschovakis. Is the Euclidean algorithm optimal among its peers? *Bulletin of Symbolic Logic* 10(3):390–418, 2004. doi:10.2178/bsl/1102022663.

[9] Lou van den Dries and Yiannis N. Moschovakis. Arithmetic complexity. *ACM Transactions on Computational Logic* 10(1), Article 2, 49 pp., 2009. doi:10.1145/1459010.1459012.